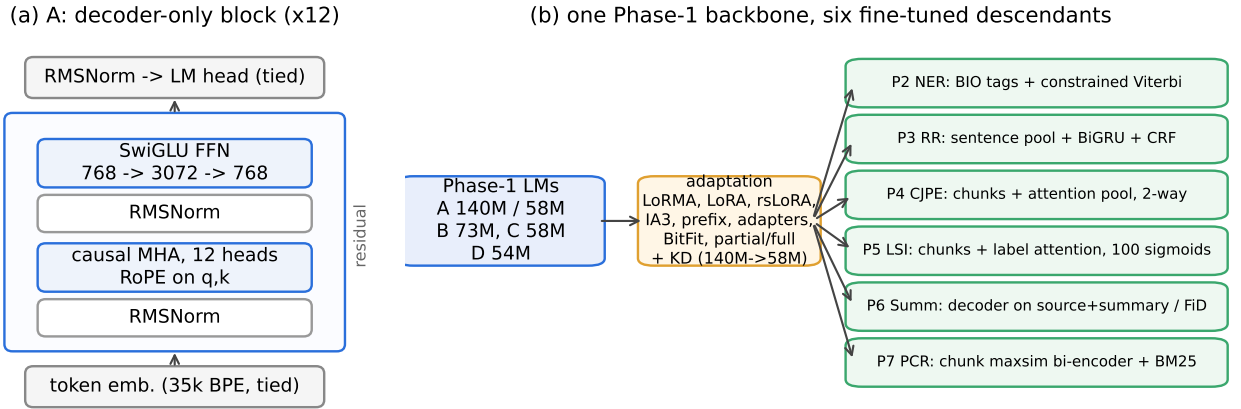
(a) A: decoder-only block (x12)
RMSNorm -> LM head (tied)
SwiGLU FFN
768 -> 3072 -> 768
RMSNorm
causal MHA, 12 heads
RoPE on q,k
RMSNorm
residual
token emb. (35k BPE, tied)
(b) one Phase-1 backbone, six fine-tuned descendants
Phase-1 LMs
A 140M / 58M
B 73M, C 58M
D 54M
adaptation
LoRMA, LoRA, rsLoRA,
IA3, prefix, adapters,
BitFit, partial/full
+ KD (140M->58M)
P2 NER: BIO tags + constrained Viterbi
P3 RR: sentence pool + BiGRU + CRF
P4 CJPE: chunks + attention pool, 2-way
P5 LSI: chunks + label attention, 100 sigmoids
P6 Summ: decoder on source+summary / FiD
P7 PCR: chunk maxsim bi-encoder + BM25